

MAIA INTELLIGENCE

Blueprint v1.1

Central Intelligence Layer for the MAIA Ecosystem

Document purpose

Define the complete product architecture, data map, security governance, memory lifecycle, model routing, cost architecture, event system, notification logic, audit layer, recovery strategy and MVP roadmap for MAIA Intelligence before development begins.

Prepared for: MAIA Business / MAIA Intelligence Core

Date: 17 June 2026

Classification: Confidential

Executive Summary

MAIA Intelligence is the single central intelligence layer that powers the MAIA ecosystem. It is not a standalone chatbot and not a separate assistant per product. The system must understand the user, tenant, product context, permissions, memory, documents, structured data and available actions before generating or executing anything.

- One MAIA Core serves Mikaelson OS, OFF, Orbit and MAIA Business through product contexts.
- The platform operates at authority levels 1 to 4: assistant, analyst, consultant and controlled operator.
- Memory follows an executive governance model: create, classify, merge, update, version, archive and delete only when necessary.
- Sensitive data is isolated, encrypted, audit-logged and protected by a six-digit Sensitive Access Code.
- The architecture includes multi-tenant isolation, model routing, event processing, notification intelligence, cost governance, AI auditability and recovery/failover.

Blueprint Section	What It Defines
Vision & Purpose	Why MAIA exists and what it is allowed to become.
System Architecture	The central core, engines, security layer and connected products.
Multi-Tenant Architecture	How data from different organizations stays isolated.
Memory Governance	How MAIA remembers, updates, merges, archives and deletes information.
Security & Sensitive Access	How confidential and sensitive data are protected.
Model Router	How MAIA can use multiple AI providers without lock-in.
Event & Notification System	How MAIA detects risk, prioritizes and alerts users.
Audit & Recovery	How MAIA explains reasoning and survives failures.
MVP & Roadmap	What to build first and what to delay.

1. Vision & Purpose

MAIA Intelligence is the central intelligence infrastructure for the MAIA ecosystem. It connects and operates across products without fragmenting into multiple disconnected AIs. Each product receives MAIA through a product context that defines tone, tools, data boundaries and operational authority.

1.1 Strategic Definition

MAIA Intelligence is the central intelligence layer that connects, understands, analyzes, recommends and executes controlled business actions across the MAIA ecosystem.


1.2 Operating Philosophy

- MAIA should never fabricate business information when sources are missing.
- MAIA should know what it is allowed to see before it attempts to answer.
- MAIA should store strategic memory, not conversational noise.
- MAIA should propose first and execute only when permission and confirmation rules allow it.
- MAIA should be designed as infrastructure, not decoration.

1.3 Product Scope

Product	MAIA Role	Initial Scope
Mikaelson OS	Executive operating intelligence.	Projects, KPIs, tasks, risks, reports, quotes.
OFF	Editorial and community intelligence.	Chapters, members, content, recommendations.
Orbit	Fiscal and invoicing intelligence.	RFC, CFDI, tickets, invoices, fiscal profile.
MAIA Business	Strategy and client intelligence.	Clients, proposals, brand, growth and operations.

2. System Architecture




2.1 Architectural Decision

MAIA must be implemented as a central service named MAIA Intelligence Core. Product applications connect to it using secure APIs. This prevents each product from building its own separate memory, tool layer and AI behavior.

Layer	Responsibility	Why It Matters
Product Apps	Mikaelson, OFF, Orbit and MAIA Business interfaces.	They provide user experience and product context.
MAIA Core	Central AI orchestration, memory, reasoning and actions.	One brain, one governance model.
Shared Data Layer	PostgreSQL, Prisma, pgvector and private storage.	Structured and semantic data live in governed systems.
Security Layer	Permissions, sensitive sessions and audit logs.	Prevents irresponsible access and data leakage.
Event Layer	Captures system signals and triggers notifications.	Lets MAIA act proactively without full autonomy.

2.2 Request Processing Flow



1. Resolve tenant, organization, user and product context.
2. Validate role and permissions.
3. Check whether sensitive access is required.
4. Retrieve relevant memory and documents.
5. Call structured tools only when needed.
6. Route the request to the right model/provider.
7. Generate reasoning output grounded in available data.
8. Create an audit trace.
9. Return answer, recommendation or action proposal.
10. Execute only if allowed, confirmed and logged.

3. Core Engines

3.1 Identity Engine

The Identity Engine resolves who is using MAIA, which tenant and organization they belong to, which product they are using and what permissions apply.

Capability	Required Behavior
Authentication	Validate session before any data access.
Tenant resolution	Every scoped request must carry tenantId.
Role resolution	Resolve admin, owner, client, member, user or service role.
Permission check	Check read/write/execute rights before memory, documents, tools and sensitive data.
Product context validation	Confirm the user can access the product context they are using.

3.2 Memory Engine

The Memory Engine uses an executive memory model. MAIA stores facts, decisions, preferences, risks and strategic context only when useful for future reasoning or execution.

3.3 Knowledge Engine

The Knowledge Engine processes real documents and sources. It extracts text, creates chunks, stores embeddings and retrieves relevant information according to permissions.

3.4 Reasoning Engine


The Reasoning Engine assembles product context, memory, document chunks, tool results and model output into a grounded answer. It must state uncertainty or missing data clearly.

3.5 Action Engine

The Action Engine turns recommendations into controlled operations. MAIA can act as assistant, analyst, consultant and operator. Autonomous director behavior is deferred to a later phase.

Authority Level	Capability	Included
1. Assistant	Answer and explain.	Yes
2. Analyst	Analyze data, documents and risks.	Yes
3. Consultant	Recommend decisions and next steps.	Yes
4. Operator	Execute approved tools with permissions and logs.	Yes, controlled
5. Autonomous Director	Execute under preapproved policies without constant confirmation.	Future

4. Multi-Tenant Architecture



MAIA must be designed for multiple organizations from the beginning. Even if the first use case is internal, the platform must prevent cross-tenant leakage before it becomes expensive to fix. Delightful how security is cheaper before disaster, apparently.


4.1 Tenant Isolation Rules

Rule	Specification
Tenant ID required	Every scoped table must include tenantId unless explicitly global.
Default deny	If tenant scope is missing, the request fails.
No cross-tenant reads	A user cannot read another tenant's data unless they are MAIA super admin with logged access.
No cross-tenant memory	Memory retrieval must filter by tenantId, ownerId, product context and security level.
No shared sensitive data	SensitiveProfile and SensitiveSession are tenant- and user-scoped.
Admin audit	Every super-admin access to tenant data must generate an audit log.

4.2 Recommended Data Isolation Model

Option	Recommendation	Reason
Shared database with tenantId	Use for MVP v1.1.	Lower complexity and faster development.
Row-level security policies	Add as platform matures.	Stronger enforcement inside database.
Dedicated database per enterprise tenant	Future enterprise option.	Useful for regulated/high-value clients.

5. Memory Governance



Memory is an asset and a liability. If unmanaged, it becomes a landfill with embeddings. MAIA needs rules for creation, update, merge, archive and deletion.

5.1 Memory Lifecycle

Stage	Behavior	Automation Level
Create	Store only useful facts, decisions, preferences, risks, strategy or sensitive data.	Semi-automatic
Classify	Assign owner, type, source, confidence and security level.	Automatic
Match	Search for existing related memories before creating duplicates.	Automatic
Merge	Combine equivalent memories and keep the clearest version.	Automatic with log
Update	If new information changes an existing fact, create a new version and mark current.	Automatic, review if sensitive
Archive	Move obsolete memories out of active retrieval while preserving history.	Automatic or admin
Delete	Remove only when required: user request, legal need, unsafe data, duplicate error.	Restricted

5.2 Memory Conflict Policy

Situation	MAIA Behavior
New memory contradicts old memory	Version the old memory, mark the new one active and log source.
Low-confidence memory	Store in review queue or ask user confirmation.
Sensitive memory detected	Require sensitive classification and confirmation before saving.
Duplicate memory detected	Merge, preserve source references and avoid repeated retrieval noise.
User asks to forget	Archive or delete according to data type and legal/security policy.

5.3 Memory Fields

Field	Purpose
id	Unique memory identifier.
tenantId	Tenant boundary.
ownerType / ownerId	User, organization, product or global ownership.
type	user_fact, client_fact, decision, preference, risk, strategy, fiscal_data, legal_data.
content	Human-readable memory.
structuredPayload	Optional structured JSON for machine use.
source	Chat, document, admin, tool, event.
confidence	Numerical confidence score.
securityLevel	Public, internal, confidential or sensitive.
version	Memory version number.
status	active, archived, deleted, review.


6. Security Architecture & Sensitive Access

Security comes before intelligence execution. MAIA must validate access before reading data, retrieving sensitive memory or executing any tool with business impact.

6.1 Security Levels

Level	Name	Examples	Controls
1	Public	Public MAIA services, public articles.	No special controls.
2	Internal	Tasks, workflows, product settings.	Authenticated access.
3	Confidential	Clients, prices, proposals, project budgets.	Role and tenant permissions.
4	Sensitive	RFC, CFDI, invoices, contracts, legal/fiscal records.	Encryption, code gate, audit, confirmation.

6.2 Sensitive Access Flow



6.3 Sensitive Access Code


Users may create a unique six-digit code. The code unlocks a temporary Sensitive Session. The code is never stored in plain text; it must be hashed using a strong password hashing approach such as bcrypt or Argon2.

Mode	What MAIA Can Do	Expiration
Locked	Know that sensitive data exists but not reveal or use full values.	Default
Sensitive Read	Read/display/summarize sensitive data allowed by role.	Short timeout, logout or manual lock
Sensitive Full	Use sensitive data during the session for approved operations.	Logout, timeout, device change or revocation

6.4 Actions Always Requiring Final Confirmation

- Issue, send or cancel invoices.
- Send fiscal data or legal documents externally.
- Modify RFC, CFDI use, fiscal regime, bank account or legal profile.
- Delete contracts, invoices or fiscal records.
- Share confidential contracts or pricing.
- Send final quotes or activate payments.

7. Data Map



7.1 Core Entities

Entity	Purpose	Important Notes
Tenant	Top-level boundary for organizations and users.	Required for multi-tenant isolation.
User	Person using MAIA.	Can belong to multiple organizations/contexts.
Organization	Company/client/community.	Owns projects, documents and memories.
ProductContext	Defines behavior and available tools.	Mikaelson, OFF, Orbit, MAIA Business.
Memory	Executive facts and decisions.	Governed lifecycle and security level.
SensitiveProfile	Fiscal/legal/personal protected profile.	Encrypted and code-gated.
Conversation	Chat thread.	Separated by user, product and tenant.
Message	Individual message.	Not automatically long-term memory.
Document	Uploaded/connected file.	Needs owner and security level.
DocumentChunk	RAG unit with embedding.	Inherits document permissions.
Project	Operational project.	Risk and budget fields.
Task	Work item.	May be created by MAIA.
Tool	Executable function.	Permission and confirmation required.
ToolLog	Record of tool execution.	Audit critical.

Event	System signal.	Used for proactive detection.
Notification	User-facing alert.	Risk, priority and recommendation.

7.2 Sensitive Entities

Entity	Fields	Rules
SensitiveAccess	userId, codeHash, failedAttempts, lockedUntil.	Code never stored plain; lock after repeated failures.
SensitiveSession	userId, tenantId, mode, deviceId, startedAt, expiresAt.	Short-lived and revocable.
SensitiveAuditLog	userId, action, resourceType, result, reason, timestamp.	Required for every sensitive access.
SensitiveProfile	legalName, RFC, CFDI, regime, zip, encryptedPayload.	Encrypted at rest.

8. Product Context Layer


MAIA stays one intelligence but responds differently depending on product context. This is behavior configuration, not separate model duplication.

Context	Tone	Default Tools	Limits
Mikaelson OS	Executive, operational.	projects, tasks, KPIs, reports, risks.	No major updates without confirmation.
OFF	Editorial, reflective, structured.	articles, chapters, members, content plans.	No access to fiscal/business sensitive data.
Orbit	Fiscal, precise, secure.	fiscal profiles, tickets, invoices, CFDI.	Sensitive code required for protected data.
MAIA Business	Strategic, premium, consultive.	clients, proposals, brand, growth, delivery.	Client confidentiality enforced.
Admin Core	Technical, audit-focused.	users, logs, memories, documents, tools.	Super-admin actions fully logged.

8.1 Product Context Payload

Field	Required	Example
tenantId	Yes	tenant_maia
productContext	Yes	mikaelson
userId	Yes	usr_123
organizationId	When applicable	org_bet
sessionId	Yes	sess_456
permissions	Yes	project.read, task.write
sensitiveSessionMode	Yes	none/read/full
locale	Recommended	es-MX

9. MAIA Model Router



The Model Router prevents MAIA from depending permanently on one AI provider. It selects the appropriate model based on task type, privacy level, cost, latency, quality and availability.

Routing Factor	Decision Logic
Task type	Reasoning, classification, extraction, summarization, code, long context or tool planning.
Sensitivity	Sensitive content may require stricter provider rules or future local/private models.
Cost	Use smaller models for classification and extraction where quality is sufficient.
Latency	Use faster models for UI interactions and heavier models for deeper analysis.
Availability	Fallback to approved alternate provider when the primary fails.
Quality tier	Strategic analysis and high-risk actions use stronger reasoning models.

9.1 Initial Provider Strategy

Task	Recommended Provider Strategy
Streaming executive chat	OpenAI primary with provider abstraction.
Long document strategy	Provider selected by context length and quality.
Classification / memory tagging	Small/cheaper model where possible.
Sensitive actions	Provider allowed only after privacy policy check.
Future private workloads	Local/private model option.

10. Cost Engine

The Cost Engine tracks and controls AI, storage, embeddings, document processing and tool execution costs. Without this layer, MAIA may scale usage faster than revenue. Humanity calls this “growth”; finance calls it “please stop.”

Cost Driver	What to Track	Control Mechanism
Chat completions	Tokens by user, tenant, product and model.	Usage quotas, caching, model routing.
Embeddings	Document chunks and memory embeddings.	Deduplicate chunks, batch embeddings.
Document storage	File size and retention.	Limits by plan and archive policy.
Vector search	Queries and retrieved chunks.	Top-k limits and retrieval optimization.
Tool execution	External API calls and retries.	Rate limits and tool budgets.
Notifications	Generated alerts and delivery channel.	Priority thresholds.

10.1 Cost Entities

Entity	Purpose
UsageRecord	Stores model calls, tokens, provider, product and tenant.
CostPolicy	Defines limits per plan, user, product or tenant.
BudgetAlert	Warns when usage crosses threshold.

ModelCostRate	Stores approximate provider pricing to estimate cost.
DocumentProcessingLog	Tracks extraction, chunking, embedding and errors.

11. Event System & Notification Intelligence



The Event System lets MAIA react to important signals without waiting for the user to ask. It should read events, analyze severity, prioritize what matters, notify the user and recommend action.

11.1 Event Types

Event Type	Example	Potential Notification
project_risk	A project deadline is near and tasks remain open.	Risk of delay detected.
payment_issue	Invoice pending or failed.	Payment/fiscal attention required.
document_uploaded	New contract uploaded.	Document ready for review.
memory_conflict	New information contradicts previous memory.	Memory update needs review.
client_activity	Client approves/rejects quote.	Recommended next step.
system_failure	Model/tool/storage failure.	System degraded or action failed.

11.2 Notification Logic

Step	Description
Detect	Receive or poll system signals.
Score	Evaluate severity, urgency and confidence.
Prioritize	Avoid notifying everything; notify what matters.
Explain	Tell the user what happened and why it matters.
Recommend	Give one or more next actions.
Track	Log notification status: sent, read, dismissed, acted.

11.3 Notification Center Requirements

- In-app notification center in MVP.
- Severity labels: info, attention, risk, critical.
- Each notification includes recommended action.
- Users can dismiss, mark done or ask MAIA for explanation.
- Later channels: email, WhatsApp, mobile push.

12. AI Audit Layer

The AI Audit Layer makes MAIA explainable. For business users, trust requires knowing why MAIA reached a conclusion, which sources it used and which tools it executed.

Audit Item	Required Detail
Input context	Tenant, product, user role and permissions used.
Memory used	Memory IDs, security level and relevance score.
Documents used	Document IDs, chunk references and retrieval score.
Tools used	Tool name, input summary, output summary and status.
Model used	Provider, model, token usage and routing reason.
Decision trace	High-level explanation of conclusion, not hidden chain-of-thought.
Action log	Whether action was proposed, confirmed, executed or denied.

12.1 User-Facing Explanation

MAIA should provide a concise explanation when asked: "Why did you recommend this?" It should reference sources, memories and tool results without exposing private system prompts or raw hidden reasoning.


Audit standard

Do not expose hidden chain-of-thought. Instead, provide a safe reasoning summary: sources used, factors considered, conclusion and recommended action.

12.2 Audit Retention & Review

Audit Record	Retention Purpose
AITrace	Debug reasoning quality, source usage and model routing.
ToolLog	Confirm what MAIA executed, when and under whose permission.
SensitiveAuditLog	Prove sensitive data access was authorized and controlled.
NotificationLog	Track whether risks were sent, read, dismissed or acted on.
IncidentLog	Analyze failures and recovery actions.

13. Recovery, Failover & Safe Mode



The Recovery System defines how MAIA behaves when something fails. The goal is not to pretend everything works. The goal is to fail safely, explain clearly and prevent unauthorized or corrupted actions.

13.1 Failure Categories

Failure Type	Example	Required Behavior
Model failure	Primary AI provider unavailable.	Retry, route to fallback model or show degraded status.
Database failure	PostgreSQL unavailable.	Stop write actions, use read-only cached state if safe.
Storage failure	Document cannot be retrieved.	Do not answer from unavailable document; explain missing source.
Tool failure	Invoice tool/API fails.	Do not mark action complete; log failure and notify user.
Embedding failure	Document cannot be indexed.	Store document pending processing and notify admin/user.
Permission failure	User lacks access.	Deny request and log attempt.
Security anomaly	Repeated code failures.	Lock sensitive access and alert user/admin.

13.2 Safe Mode

Safe Mode is activated during critical failures or security uncertainty. In Safe Mode, MAIA may answer general questions but cannot execute sensitive or irreversible actions.

- No sensitive data access unless security systems are fully available.
- No invoice issuing, deletion, external sending or legal/fiscal modification.
- Use cached data only if marked safe and current enough.
- Clearly label responses as degraded when information may be incomplete.
- Create incident logs for admin review.

13.3 Recovery Standards

Standard	Specification
Retry policy	Use limited retries with exponential backoff for transient failures.
Fallback policy	Use approved alternate providers/tools only when privacy and permissions allow.
Incident logging	Every failure must capture source, severity, user impact and resolution status.
User communication	Never fake success; tell the user when something failed.
Post-recovery validation	Verify failed jobs before resuming normal operations.

14. Technical Architecture

Layer	Recommendation	Reason
Frontend	Next.js + TypeScript	Fast product UI with strong ecosystem.
Backend	Next.js API routes initially; modular service architecture.	Efficient MVP, scalable later.
Database	PostgreSQL / Neon	Structured data and reliable relational model.
ORM	Prisma	Schema control and migrations.
Vector Search	pgvector	Embeddings inside PostgreSQL.
AI Layer	Provider abstraction + Model Router	Avoid lock-in and optimize cost/quality.
Storage	Private object storage	Secure document handling outside DB.
Auth	Secure auth with sessions and roles	Required for tenant and sensitive access.
Jobs	Background queue for documents/events	Needed for processing and notifications.
Observability	Logs, metrics and error tracking	Needed for recovery and audit.

14.1 Repository Structure

```

maia-intelligence/
  src/
    app/
    components/
    core/
      identity/
      tenancy/
      memory/
      knowledge/
      reasoning/
      actions/
      events/
      notifications/
      cost/
      model-router/
    products/
      mikaelson/
      off/
      orbit/
      maia-business/
    security/
      permissions/
      sensitive-access/
      audit/
      recovery/
    integrations/
    lib/
      ai/
      db/
      storage/
      jobs/
    dashboard/
    prisma/
      schema.prisma
      migrations/
    docs/
    tests/

```

15. MVP v1.1 Definition

15.1 Included

- MAIA Core service.
- Multi-tenant identity and permission foundation.
- Product Context Layer.
- Executive Memory Engine with lifecycle governance.
- Document upload, extraction, chunking, embeddings and RAG.
- Reasoning Engine with grounded responses.
- MAIA Model Router abstraction.
- Sensitive Access Code and Sensitive Session.

- Basic Action Engine: create task, save memory, generate report, draft response, create quote draft.
- Tool logs, sensitive logs and AI audit trace.
- Event System and in-app Notification Center.
- Cost Engine basic usage tracking.
- Recovery/Safe Mode policies.
- Admin dashboard for users, tenants, memories, documents, logs, tools, notifications and usage.

15.2 Not Included Yet

- Voice assistant.
- WhatsApp integration.
- Live Gmail/Calendar integrations.
- Stripe payment automation.
- Fully autonomous agents / Level 5.
- Automatic invoice issuing without human confirmation.
- Complex multi-agent workflows.
- Enterprise billing and plan enforcement beyond basic cost tracking.

15.3 First Integration

Mikaelson OS should be the first product integration because it is closest to MAIA's operational value: projects, KPIs, risks, reports, tasks, executive decisions and controlled actions.

16. Development Roadmap



Sprint	Objective	Deliverables
1	Core + tenancy	Auth, tenant, org, user, roles, product contexts.
2	Memory governance	Memory schema, classification, merge, update, version, archive.
3	Knowledge + RAG	Document upload, extraction, chunks, embeddings, retrieval.
4	Router + reasoning	Model Router, context assembly, streaming chat, grounded responses.
5	Action + audit	Tool registry, task/report tools, confirmations, audit traces.
6	Events + notifications	Event bus, risk scoring, notification center, recommendations.
7	Mikaelson integration	Projects, tasks, KPIs, risk alerts and executive dashboard integration.
Later	Advanced operations	Voice, email/calendar, payments, WhatsApp, autonomous agents.

17. Acceptance Criteria

Area	Acceptance Criteria
Tenant isolation	No cross-tenant access without super-admin audit.
Memory	Classify, retrieve, update, merge and archive executive memory.
Sensitive access	Code-gated session plus audit logs.
Documents	Permission-aware RAG from uploaded documents.
Model Router	Provider abstraction with usage tracking.
Cost Engine	Usage logged by tenant, product and model.
Events	Prioritized risk and operational notifications.
AI Audit	Safe explanation of sources, tools and model used.
Recovery	Safe mode or clear failure message, never fake success.
Actions	Permission, confirmation and log for execution.

18. Final Product Positioning

MAIA Intelligence should be positioned as the central operating intelligence layer of the MAIA ecosystem. The first release must prove four things: MAIA remembers strategically, reads real information, protects sensitive data and operates controlled actions. Later releases can add live integrations, voice and autonomous agents.

Strategic standard

Build the brain before adding limbs. MAIA Core must be stable before connecting every external tool humans have invented to make their calendars look like crime scenes.

18.1 Final Build Rules

Rule	Meaning
One brain	Do not create separate AI systems per product. Use one MAIA Core with product contexts.
Data before action	MAIA must validate sources and permissions before answering or executing.
Memory with governance	Every memory must be classified, versioned and auditable.
Security by default	Sensitive data stays locked unless access is explicitly verified.
Explainable intelligence	Users must be able to see sources, tools and safe reasoning summaries.
Controlled execution	Every high-impact action requires permission, confirmation and logs.

18.2 Engineering Handoff Boundary

MAIA Intelligence Blueprint v1.1 defines the product architecture, governance model, operating philosophy and strategic build direction of MAIA Intelligence.

This document intentionally does not define implementation details such as database schemas, Prisma models, API contracts, API payloads, permission matrices, threat models, infrastructure deployment standards or engineering procedures.

Those engineering specifications belong exclusively to MAIA Technical Architecture v1.

The Blueprint is considered complete when it clearly defines:

- What MAIA is.
- What MAIA is allowed to do.
- What data MAIA governs.
- How MAIA protects sensitive information.
- What capabilities belong to the MVP.
- What must be built before advanced automation.

Boundary Area		Blueprint Owns	Technical Architecture Owns
Purpose	Vision, strategic direction, ecosystem scope and operating philosophy.		Engineering implementation, system design and deployment architecture.
Authority Model	Levels 1–5 authority framework, execution philosophy and governance rules.		Technical enforcement of permissions, execution controls and approval flows.
Product Scope	Product definitions, ecosystem boundaries, MVP scope and roadmap.		Sprint planning, implementation milestones and engineering deliverables.
Product Contexts	Context definitions, behavioral boundaries and product responsibilities.		Context payloads, routing logic and technical context enforcement.
Memory	Governance model, lifecycle principles, classification philosophy and retention intent.		Memory scoring, retrieval logic, storage architecture, versioning and indexing.
Security	Security philosophy, sensitive access model, confirmation standards and governance rules.		Threat model, permission matrix, encryption, tenant isolation and endpoint security.
Actions	Controlled execution philosophy, authority boundaries and approval requirements.		Tool registry, idempotency, execution controls, audit enforcement and confirmation flows.
Events & Notifications	Event concepts, notification objectives and operational intent.		Event processing architecture, notification scoring, prioritization and delivery logic.
Cost Governance	Cost control philosophy, usage accountability and budget governance.		Cost engine implementation, usage tracking, limits, quotas and budget enforcement.
Audit & Explainability	Transparency requirements, trust model and explainability standards.		Audit records, traces, observability, monitoring and compliance logging.
Recovery & Continuity	Recovery philosophy, safe-mode principles and resilience expectations.		Backup strategy, disaster recovery, failover mechanisms and recovery procedures.
Readiness	Strategic approval to build MAIA Core and validate product assumptions.		Technical approval to develop, test, deploy and scale the platform.

Blueprint Status: APPROVED as the strategic and governance source of truth for MAIA Intelligence.

Technical Architecture Status: APPROVED as the engineering source of truth for implementation, deployment, operations and scaling.

MAKE IT UNSTOPPABLE